

CNNs In Depth

Lamia Zain lamiahasan4@gmail.com LinkedIn

June,21, 2024

## Connect Sessions | Purpose

#### A Connect Session IS:

- Focused on learning, encouragement & graduation for a group of students coached by a Udacity Session Lead
- Setting weekly study goals
- Helping each other with progress (including peer to peer)
- Keeping everyone accountable for their responsibilities
- A way to meet individuals in tech field & learn about the industry
- Mandatory

#### A Connect Session IS NOT:

- A social meetup
- A study group
- A substitute for online learning
- Optional





# Let's check your progress

You are encouraged to spend at lest 10 hours/week to graduate.



Presentation date

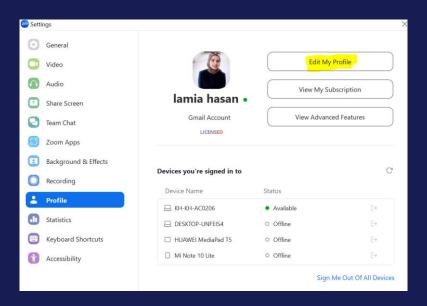
## U UDACITY

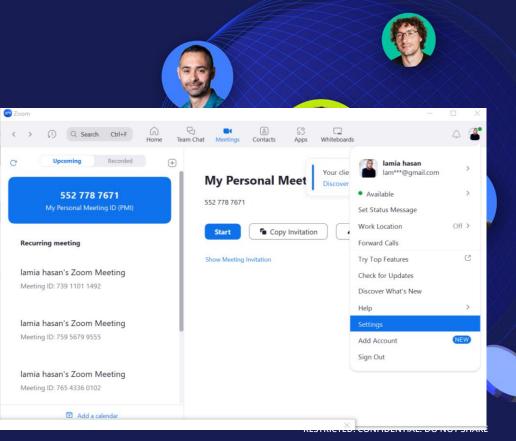
# Attendance is taken automatically

Please change your name to be First Name and Last name on Zoom Like: Lamia Zain



# **Change your**Name on Zoom





# **UDACITY** Change your Name on Zoom

Products

Solutions

Resources

Personal

zoom

Profile

Meetings

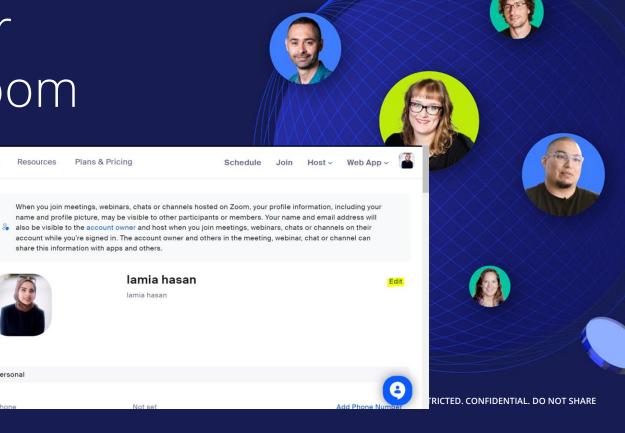
Webinars

Personal Contacts Personal Devices

Whiteboards

Surveys NEW Recordings Scheduler

Settings Reports



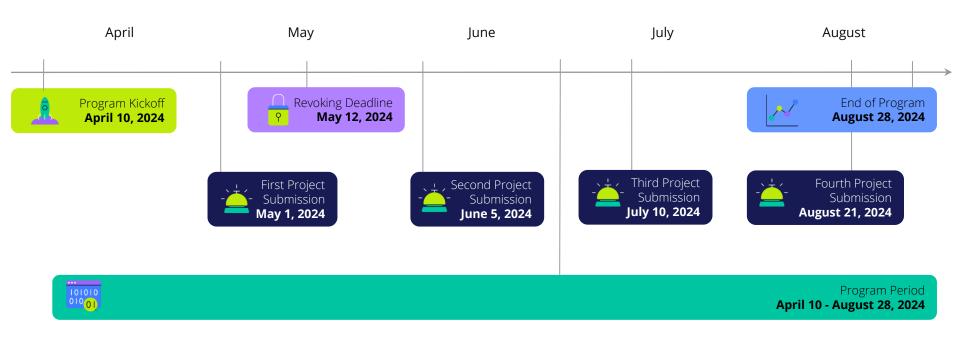
# Session Lead role:

#### **Communication Chart**

Issue	Where to go?
Classroom access/ Withdrawal/ Graduation issues/ Plagiarism/ Project Review Inquiries	Email support@udacity.com
Technical Issues, Attendance, Content Related Issues/ Project inquiries	Session Lead
Session Switch/ Community related issues	Community Moderators



2024





## Four-weeks Agenda, Weekly schedule

Week 10	Jun 12, 2024			Finish the lessons below from the Convolutional Neural Networks Introduction to CNNs CNN Concepts  [Work on/submit the #3 project: Landmark Classification & Tagging for Social Medial	Convolutional Neural Networks Introduction to CNNs CNN Concepts
Week 11	Jun 19, 2024			Finish the lessons below from the Convolutional Neural Networks CNNs in Depth  [Work on/submit the #3 project: Landmark Classification & Tagging for Social Media]	Convolutional Neural Networks CNNs in Depth
Week 12	Jun 26, 2024			Finish the lessons below from the Convolutional Neural Networks Transfer Learning [Work on/submit the #3 project: Landmark Classification & Tagging for Social Media]	Convolutional Neural Networks Transfer Learning
Week 13	Jul 3, 2024			Finish the lessons below from the Convolutional Neural Networks Autoencoders  [Work on/submit the #3 project: Landmark Classification & Tagging for Social Media]	Convolutional Neural Networks Autoencoders Project Walkthrough: Landmark Classification & Tagging for Social Media
Week 14	Jul 10, 2024	Jul 10, 2024	Landmark Classification & Tagging for Social Media	Finish the lessons below from the Convolutional Neural Networks Object Detection and Segmentation  [Work on/submit the #3 project: Landmark Classification & Tagging for Social Media]	Convolutional Neural Networks Object Detection and Segmentation Project Walkthrough: Landmark Classification & Tagging for Social Media



### Student Milestone | Revoking

#### **REVOKING**

**Revoking** is the process by which Udacity removes a student from a Nanodegree program.

AWS reserves the right to revoke you from the program if you do not comply with program requirements.

#### **CRITERIA**

Students can be revoked if they fail to:

- Submit Project 1
- Complete the required concepts







## Code of Conduct | Plagiarism

#### **BASIC RULES**

- Project submissions must consist of original work
- Submitted projects will be scanned for plagiarism
- Students who are found to have plagiarised will risk their Nanodegree being revoked
- Read the honor code and the rubric carefully for all projects









# Recap



The difference between CNNS and multi perceptron networks when classifying Images

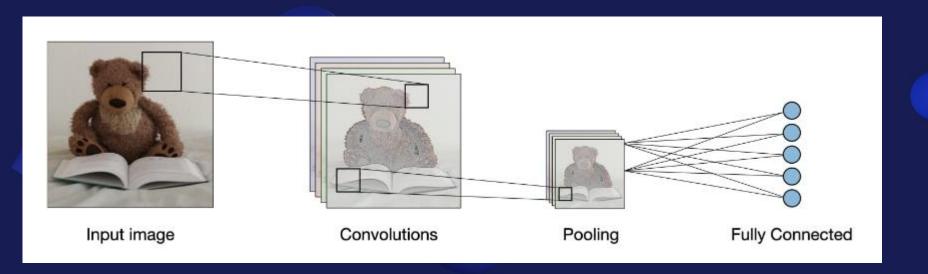


	CNNs	MLPs
Architecture	<ul> <li>Convolutional layers that apply filters to the input image. Capturing local spatial patterns.</li> <li>followed by pooling layers that downsample the feature maps</li> </ul>	<ul> <li>MLPs consist of fully connected layers where each neuron is connected to every neuron in the previous and subsequent layers</li> <li>MLPs do not take into account the spatial structure of the input data.</li> <li>Too Unnecessary Many Parameters</li> </ul>
Parameter Sharing	<ul> <li>Use parameter sharing, i.e, the same filter is applied to different regions of the input image.</li> <li>Smaller number of learned parameters compared to MLPs</li> <li>Hence better for huge input data like images</li> </ul>	<ul> <li>Each neuron is trained independently</li> <li>Network learns unique weights for every connection.</li> <li>Large number of parameters.</li> <li>Computationally expensive for image data.</li> </ul>
Translation Invariance	Due to their use of convolutional and pooling layers, they can detect local patterns regardless of their position in the image.	Sensitive to the position of pixels in the input image, making them less robust to translation or slight changes in input position.

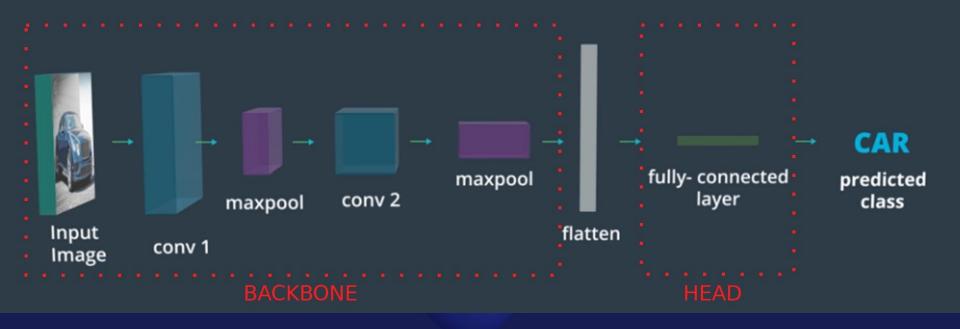




## <u>Cs230 – Stanford University</u>







#### **Convolution Layer**

- Responsible for preserving special information about an image.
- Detects features of images.
- By applying learnable Kernels/filters (Wights in MLPs).
- To produce feature maps.
- Size of the output feature maps is dependent on the padding and stride.

#### **Example**

- Input image of size 8x8.
- Kernel size of 3x3 stride of 1
- Padding= "valid"

-1	0	1
-2	0	2
-1	0	1
	Kerne	7

_								
	0.1	0.2	0.3	0.4	0.5	0.6	0.7	0.8
	0.9	0.8	0.7	0.6	0.5	0.4	0.3	0.2
	0.1	0.2	0.3	0.4	0.5	0.6	0.7	0.8
	0.9	0.8	0.7	0.6	0.5	0.4	0.3	0.2
	0.1	0.2	0.3	0.4	0.5	0.6	0.7	0.8
	0.9	0.8	0.7	0.6	0.5	0.4	0.3	0.2
	0.1	0.2	0.3	0.4	0.5	0.6	0.7	0.8
	0.9	0.8	0.7	0.6	0.5	0.4	0.3	0.2

0.1	0.2	0.3	0.4	0.5	0.6	0.7	8.0
0.9	8.0	0.7	0.6	0.5	0.4	0.3	0.2
0.1	0.2	0.3	0.4	0.5	0.4	0.7	8.0
0.9	8.0	0.7	0.6	0.5	0.6	0.3	0.2
0.1	0.2	0.3	0.4	0.5	0.4	0.7	8.0
0.9	8.0	0.7	0.6	0.5	0.6	0.3	0.2
0.1	0.2	0.3	0.4	0.5	0.4	0.7	8.0
0.9	8.0	0.7	0.6	0.5	0.6	0.3	0.2

Applying kernels

0	0	0		

Image

0	0	0	-0.2	4.6	4.9
-0.2	-0.2	-0.2	-0.4	4.2	3.9
0.2	0.2	0.2	0	4.6	4.9
0	0	0	0	4.4	4.1
0	0	0	0	4.6	4.9
0	0	0	0	4.4	4.1

Final Image



#### **Feature map Size**

$$Output \ size = \frac{Input \ size - Filter \ size + (Number \ of \ Bads \ in \ a \ direction)}{Stride} + 1$$

- <u>Input size:</u> The spatial dimensions (width or height) of the input Image.
- <u>Filter size:</u> The spatial dimensions (width or height) of the convolutional filter (kernel).
- <u>Padding:</u> The number of pixels added to the borders of the input feature map before applying the convolution. Padding helps preserve spatial dimensions and information.
- <u>Stride:</u> The number of steps the filter moves (horizontally or vertically) as it slides over the input feature map during the convolution operation.
- For the previous example,

Output Height = 
$$\frac{8-3}{1} + 1 = 6$$
  
Output width =  $\frac{8-3}{1} + 1 = 6$ 

• It's important to note that for a 2D convolution, the formula applies independently to both the width and height dimensions of the input.

#### **Convolution Layer**

- Responsible for preserving special information about an image.
- Detects features of images.
- By applying learnable Kernels/filters (Wights in MLPs).
- To produce feature maps.
- Size of the output feature maps is dependent on the padding and stride.

#### **Example**

- Input image of size 8x8.
- Kernel size of 3x3 stride of 1
- Padding= "Same"
- Padding mode='zeros'

#### **Convolution Process**

- Center the filter
- Element-wise Multiplication.
- Summation
- Replace Pixel value
- Shift and Repeat:

-1	0	1
-2	0	2
-1	0	1

#### Kernel

1.2	0.2	0.2	0.2	0.2	0.2	0.2	-1.7
2	0	0	0	-0.2	0.2	0.2	-2
2	0	0	0	-0.2	0	0.2	-2
2	0	0	0	0	0	0	-2
2	0	0	0	0	0	0	-2
2	0	0	0	0	0	0	-2
2	0	0	0	0	0	0	-2
1.8	-0.2	-0.2	-0.2	0	-0.2	-0.4	-1.3

0	0	0	0	0	0	0	0	0	0
0	0.1	0.2	0.3	0.4	0.5	0.6	0.7	0.8	0
0	0.9	0.8	0.7	0.6	0.5	0.4	0.3	0.2	0
0	0.1	0.2	0.3	0.4	0.5	0.6	0.7	0.8	0
0	0.9	0.8	0.7	0.6	0.5	0.4	0.3	0.2	0
0	0.1	0.2	0.3	0.4	0.5	0.6	0.7	0.8	0
0	0.9	0.8	0.7	0.6	0.5	0.4	0.3	0.2	0
0	0.1	0.2	0.3	0.4	0.5	0.6	0.7	0.8	0
0	0.9	0.8	0.7	0.6	0.5	0.4	0.3	0.2	0
0	0	0	0	0	0	0	0	0	0

2D image

#### **Feature map Size**

$$Output \ size = \frac{Input \ size - Filter \ size + (Number \ of \ Bads \ in \ a \ direction)}{Stride} + 1$$

- <u>Input size:</u> The spatial dimensions (width or height) of the input Image.
- <u>Filter size:</u> The spatial dimensions (width or height) of the convolutional filter (kernel).
- <u>Padding:</u> The number of pixels added to the borders of the input feature map before applying the convolution. Padding helps preserve spatial dimensions and information.
- <u>Stride</u>: The number of steps the filter moves (horizontally or vertically) as it slides over the input feature map during the convolution operation.
- For the previous example,

Output Height = 
$$\frac{8-3+2}{1} + 1 = 8$$
  
Output width =  $\frac{8-3+2}{1} + 1 = 8$ 

• It's important to note that for a 2D convolution, the formula applies independently to both the width and height dimensions of the input.

#### Stride:

Number of steps the convolutional filter (kernel) takes as it slides over the input data (Stride=1)

#### **Padding:**

- Adding extra pixels around the borders of the input data before applying the convolution operation
- Without padding, the size of the feature maps reduces with each convolution operation.
- which can lead to a progressive loss of spatial information
- In PyTorch Padding =
- "valid" (no padding),
- "same" (padding to keep output size same as input size with stride 1.

```
CONV2D &

CLASS torch.nn.Conv2d(in_channels, out_channels, kernel_size, stride=1, padding=0, dilation=1, groups=1, bias=True, padding_mode='zeros', device=None, dtype=None) [SOURCE]
```

Conv2d in Pytorch (captures spatial hierarchies in the data)

Conv1d in Pytorch (captures temporal or sequential patterns in the data)



#### **Pooling:**

- Downsamples the spatial dimensions of the feature maps obtained after convolutional layers.
- Reduces the number of parameters in the network.
- The two most common types of pooling layers are MaxPooling and AveragePooling

#### **MaxPooling:**

- The most common type of pooling used in CNNs.
- It operates on small regions (usually 2x2 or 3x3) of the input feature maps.
- Outputs the maximum value of the elements within that region.
- The stride parameter determines the step size at which the pooling window moves over the feature map.
- Typically, the stride is set to the same value as the pooling window size to avoid overlapping

CLASS torch.nn.MaxPool2d(kernel\_size, stride=None, padding=0, dilation=1, return\_indices=False, ceil\_mode=False) [SOURCE]

If you have a feature map of size 8x8 and MaxPolling layer of size (2,2) what do you think the size of the output will be ??



#### **Pooling:**

- Downsamples the spatial dimensions of the feature maps obtained after convolutional layers.
- Reduces the number of parameters in the network.
- The two most common types of pooling layers are MaxPooling and AveragePooling

#### **MaxPooling:**

1.2	0.2	0.2	0.2	0.2	0.2	0.2	-1.7
2	0	0	0	-0.2	0.2	0.2	-2
2	0	0	0	-0.2	0	0.2	-2
2	0	0	0	0	0	0	-2
2	0	0	0	0	0	0	-2
2	0	0	0	0	0	0	-2
2	0	0	0	0	0	0	-2
1.8	-0.2	-0.2	-0.2	0	-0.2	-0.4	-1.3

2	0.2	0.2	0.2
2	0	0	0.2
2	0	0	0
1.8	0	0	0

After Pooling

Output Feature Map Before Booling

#### **Pooling:**

- Downsamples the spatial dimensions of the feature maps obtained after convolutional layers.
- Reduces the number of parameters in the network.
- The two most common types of pooling layers are MaxPooling and AveragePooling

#### **AveragePooling:**

- Less Commonly used
- Computes the average of the elements within each pooling window instead of taking the maximum.

```
CLASS torch.nn.AvgPool2d(kernel_size, stride=None, padding=0, ceil_mode=False, count_include_pad=True, divisor_override=None)
```

# Break (10 minutes)

**Satisfaction Survey** 



# Train a Simple CNN for Fashion MNIST Dataset

# Why Do you think CNN Gives Less Accuracy than MLPs



Any Question?

# Thank you

